No Resolvable Agreement Between Descriptor Distance and Colorimetric Array Response: A Reanalysis of Authentic Chemical Warfare Agent Measurements

Ku Kang,[†] Jeongyun Kim,[†] Miae Kim,[†] Doo-Hee Lee,[‡] Jin Yoo,[†] and Yoon Jeong Jang[*,†]

†*Chemical, Biological, Radiological and Nuclear (CBRN) Defense Research Institute, Ministry of National Defense, Seoul, Republic of Korea*

‡*Department of Civil Engineering and Environmental Sciences, Korea Military Academy, Seoul 01805, Republic of Korea*

E-mail: yjjangchem@korea.kr

Abstract

Chemical warfare agent detection is developed and qualified on simulants, and we asked what measurement establishes that one reproduces its agent. We reanalyse the colorimetric tables of three published runs of one 29-dye array, which measured sixteen authentic agents at four concentrations in two optical channels, and compare the response patterns with the molecular-descriptor distances on which simulant selection rests. All descriptor calculations and tests are new. The array discriminates on average: in all six run-channel cells, every compound agrees with itself across concentrations more closely than with the other compounds of that cell, and all six clear $q = 0.05$ against a null that permutes compound labels within each concentration. Descriptor distance

does not track response distance at any resolvable effect size. Twelve exact Mantel tests and six combinations of them leave nothing below $q = 0.05$, against a design that resolves a rank correlation of 0.85 at 80% power at the corrected threshold the conclusions use, and 0.67 at a nominal α, where the largest observed is 0.57. That null holds under every analyst choice we enumerated, from the descriptor encoding to the panel composition, while the count of valid matches is moved by all five and is not a measurement. Measured vapour pressure, tested separately, reaches $q = 0.0430$ in the reflectance channel within its own prespecified family of two, uncorrected against the twelve descriptor cells, yet returns one valid match at $p = 0.4647$, on a gap smaller than the temperature mismatch in its own source data. Dimethyl methylphosphonate, the standard G-agent simulant, is anti-correlated with soman and cyclosarin in the ultraviolet. Across the three runs, one designated simulant has been measured beside the agents.

Introduction

Nearly every chemical warfare agent detector in the literature is developed, optimised and qualified on compounds that are not chemical warfare agents.[1,2] Dimethyl methylphosphonate stands in for the G-series nerve agents,[3] dipropylene glycol monomethyl ether for training detectors,[4] and methyl salicylate for vapour penetration through protective ensembles.[5,6] The substitution is unavoidable, because the agents are lethal at concentrations where handling them requires a licensed facility, and it is load-bearing: a sensor reported as selective for sarin has, in most cases, never been exposed to sarin.

What licenses the substitution is a claim of similarity. The compound chosen is said to be agent-like, and similarity is read from molecular descriptors: vapour pressure, dipole moment, frontier orbital energies, substructure and shape. We test the first of those separately, on measured vapour pressures, in Section . Cheminformatic selection on exactly those grounds has been proposed as the systematic way to find a simulant,[7] reviews catalogue the

compounds in use for environmental behaviour,[8] for emergency decontamination trials[9] and for toxic industrial chemicals,[10] and the US Army maintains a test operations procedure governing simulant selection for laboratory, chamber and field testing.[11] The claim that matters for a detector is different. It is that the substitute produces the response the agent would produce in the same device. Those two claims are connected only if descriptor similarity transfers to response similarity.

That transfer has been tested once, for a different measurement and in simulation. Agrawal and co-workers asked how effectively four common simulants, dimethyl methylphosphonate and diisopropyl fluorophosphate among them, mimic the adsorption properties of sarin and soman, by computing dilute-limit adsorption for all six compounds in 2969 metal–organic frameworks of known structure under two independent force fields.[12] That design answers the question for adsorption, which is the quantity a capture material is judged on. It cannot answer it for detection, because the response of a detector is not an adsorption isotherm, and because a simulation of the agent is not a measurement of it. What has not been done is the measured comparison: the agents themselves, in a detector, against the descriptor distances on which the substitution rests.

It has not been measured because measuring it requires the agents. Response data on authentic agents exist in few places, and the compounds used as simulants are rarely run alongside them. A recent survey of threat simulants records that they have generally been selected on the basis of their historical use, and reports a draft ASTM guide for explosive and contraband simulants; it describes no equivalent document on the chemical warfare agent side.[13] A critical review of the low-toxicity simulants in current use states the selection criterion plainly, as the alignment of physicochemical properties such as hydrophobicity, volatility and solubility with those of the agent, and assesses dipropylene glycol monomethyl ether and malathion among others on that basis.[14] That criterion is a statement about the compound. It is not a statement about the detector, and nothing in the review measures whether a compound selected that way produces the response the agent produces.

Asking whether the components a detection claim rests on have themselves been validated is a question this field has begun to put in print: a recent guide to choosing and validating aptamers for biosensors opens on the observation that many reported aptamers remain insufficiently characterised, and that the resulting reproducibility problem makes a reliable choice hard.[15] Simulants are the analyte-side form of that problem, and the concern itself is not new. It has been stated in this journal. Reporting a bioelectronic nose tested on authentic sarin and VX, two of us and our co-workers noted that detection research built on structural simulants may not carry over to the agents, depending on the selectivity of the sensor, and used combinatorial pattern recognition across a multi-receptor array to separate the agents themselves.[16] That work asserted the risk and worked around it by measuring the agents. It did not quantify the risk, because quantifying it requires putting the descriptor distances and the measured responses side by side and asking whether one predicts the other. The same holds for our group's field-deployable detection work, which likewise measured the agents rather than testing the substitution.[17] The contribution here is that measurement and not the assertion. Our institute is an OPCW designated laboratory, which identifies and assays authentic agents under an accredited procedure, so the comparison can be made on the agents themselves.

This work reports that measurement. A 29-dye colorimetric array was used to record the responses of sixteen authentic agents at four exposure concentrations in two optical channels, and those responses were compared against the molecular-descriptor distances that simulant selection relies on. The array separates compounds in every run and channel, so the measurement is sound. Descriptor distance nonetheless carries no detectable information about response distance. Dimethyl methylphosphonate, the compound the field treats as the G-agent surrogate, is anti-correlated with two of the four G agents, and its resemblance to sarin weakens as the concentration falls towards the range where point detectors operate. Where the designation record and the regulatory toxicity axis might have supplied the missing evidence, the record names a measurement that is not detection and the axis the agents

are placed on does not reach the incumbents.

This is not a new sensing method, and we propose no replacement compound. It is a measurement of whether the substitutes the field already relies on carry the property they are relied upon for.

Results and Discussion

The array separates the agents

A comparison between descriptor space and response space is worth making only if the response measurement resolves the compounds at all. That question is already answered in the reports these data come from, and the answer does not depend on anything we do here. On the first run a linear discriminant classifier built on a six-dye subset under dual illumination separated the four G agents, VX and dimethyl methylphosphonate at 100% accuracy, against 87.5% under visible illumination alone.[18] On the second, a six-dye subset classified all seven agents at an accuracy of 0.964 with a macro F_1 of 0.964 under five-fold stratified cross-validation.[19] On the third, a four-dye subset separated the four Novichok agents at an accuracy of 1.000 and a macro F_1 of 1.000 under leave-one-out cross-validation, over a working range of 10 to 500 μM with a detection limit near 10 μM.[20]

Three qualifications belong with those figures, and together they move where the premise rests. Each is obtained on a dye subset chosen for that run, not on the 29-dye array we analyse. The first run reports no held-out evaluation, so its 100% is an apparent rather than a validated accuracy. And the run that carries the comparison this paper is about is that first run, because it is the only one in which a designated simulant was measured beside the agents; the two cross-validated figures come from the second and third runs, which contain no simulant. A validated accuracy on a run that holds no simulant does not establish the premise for a run that does. We also cannot reproduce either cross-validated figure from the deposited tables, which are averaged over replicates, so we quote them and do not verify

them. The premise for the first run therefore rests on our own check below, which covers all six run-channel cells including that run's two.

What our own check adds is narrower and is reported because the comparison that follows is made on patterns rather than on class labels. Each exposure gives a 29-dye ΔE_{00} vector, which we centred dye-wise at each concentration and normalised to unit length, so that agreement is measured between patterns and not between magnitudes, and took agreement as the cosine. The question it answers is whether a compound's pattern is stable enough across dose to be treated as that compound's pattern at all.

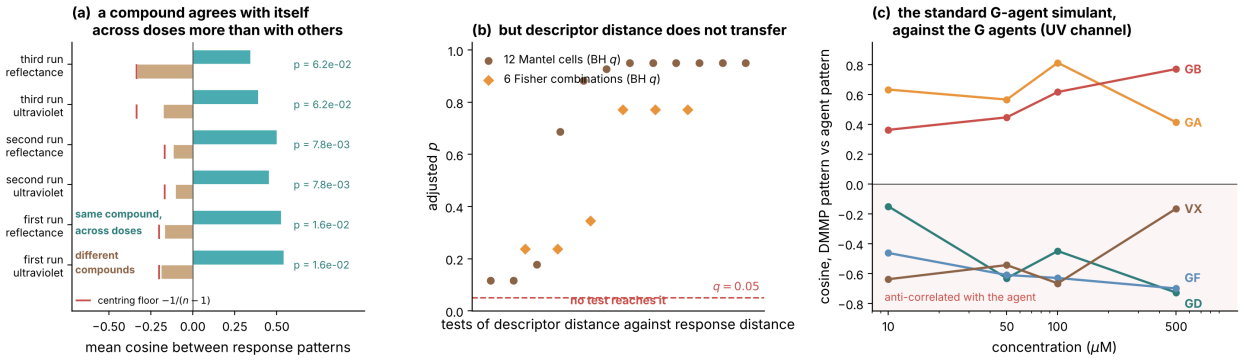


Figure 1: Response-space tests on authentic agents. (a) Mean cosine between response patterns, within one compound across concentrations and between different compounds, for three runs in two optical channels. (b) Adjusted p for every test of descriptor distance against response distance: twelve exact Mantel cells and six Fisher combinations of the independent runs. (c) Cosine between the response pattern of DMMP and those of the four G agents and VX, by concentration, in the ultraviolet channel.

It is stable enough (Figure 1a). One compound's patterns at two concentrations correlate at $+0.345$ to $+0.542$, different compounds anti-correlate at -0.103 to -0.330, and the ordering holds in all six run-channel cells. Because the independent unit is the compound and not the pair, we pre-specified a sign-flip enumeration over the four to seven compounds of each cell. That test returns the smallest value its compound count allows in all six cells, from 0.0078 to 0.0625, so what it reports is a property of the design and not of the data: centring leaves the per-compound difference positive for almost any input that is not degenerate. We withdraw it. The same statistic is tested here against a null that permutes compound

labels within each concentration, which leaves every centred vector untouched and changes only which response is assigned to which compound. All six cells clear $q = 0.05$. Four of the six are sampled at 10,000 relabellings and return the smallest value that sampler can give, 9.999×10^{-5}, so those four are censored at the sampler's resolution and the smallest adjusted value over the six, 0.00015, is an upper bound rather than an estimate. The two four-compound cells are exact, by enumeration of all 13,824 relabellings, and give one-sided values of 0.0046 and 0.0033. The effect is large enough that the censoring does not bear on the conclusion: the observed statistics run from +0.558 to +0.729 against null means of +0.102 to +0.314. Cohen's d_z over compounds runs from +0.92 to +5.12 as a descriptive effect size. Supporting Information Section S12 gives the per-cell values and why the bound the earlier version used does not apply.

The tables analysed here hold one value per compound, concentration, dye and channel, so the within-compound term is agreement of a compound with itself at a different dose and not agreement between repeated measurements. Repeats were reported in the source work for the second and third runs, for selected dye-analyte combinations in the second[19] and for all experiments in the third,[20] but the repeat-level values are not in these tables, so no repeatability estimate enters any analysis below.

Two features of the measurement belong here rather than in the Supporting Information. The first is that dye-wise centring drives the mean between-compound cosine negative by itself. An earlier version of this section compared the observed means against $-1/(n-1)$ and read the margin above it as evidence. That comparison does not hold, and Supporting Information Section S12 gives the two reasons: the centring step renormalises each vector to unit length, which destroys the zero-sum property the identity requires, and for vectors that do sum to zero $-1/(n-1)$ is an upper bound on the norm-weighted mean cosine rather than a floor. What does quantify the inflation is the permutation null itself. Relabelling compounds within each concentration leaves a mean within-minus-between difference of +0.1024 to +0.3138 across the six cells, against observed differences of +0.5578 to +0.7288. The sep-

aration exceeds what centring alone produces in every cell, including the reflectance channel of the third run that the earlier comparison set aside. What the test establishes is the per-compound statement that a compound agrees with itself across concentrations more than it agrees with the other compounds in its cell on average. It does not establish the stronger pairwise statement: in 13 of the 34 compound-and-cell instances, some other compound's consensus pattern is closer to a compound than that compound's own agreement across concentrations. The classification results reported for these runs in the source work, not this cosine, are what establish that the agents are separable, and we quote those rather than reproduce them, because the deposited tables are averaged over replicates and the cross-validation design cannot be rebuilt from them. The second is that six dye columns in the ultraviolet pivot of the second run were bit-identical across every row. Tracing them to the source images showed that those dyes do not fluoresce under ultraviolet excitation and the regions of interest had read plate background, so one channel had been recorded six times. Columns identical across all rows are counted once, which leaves 23 effective channels in that cell and 25 in the ultraviolet pivot of the third run. Supporting Information Section S21 sets out what the array does and does not establish in more detail.

Descriptor distance does not predict response distance

Simulant selection rests on an implicit transfer: that a compound close to an agent in descriptor space will be close to it in whatever the detector measures. We tested the transfer directly, as the agreement between two distance matrices over the same compounds, one built from descriptors and one from response patterns. Agreement was measured by Mantel correlation and referred to a null built by enumerating every relabelling of the compounds, so the test is exact rather than sampled.

The test is run within each run and never across runs. The three runs share no compound, so the dye-wise sensitivity of one plate cannot be calibrated against another, and pooling responses across them would require a commutability assumption we cannot support. That

restriction sets the power. Six compounds in the first run and six in the second carry both a response pattern and a descriptor vector, giving fifteen pairs in each, and the third run drops out because the A-series compounds are not in the descriptor panel, which follows from the source itself: it reports the four as the codes A-230, A-232, A-234 and A-242, gives their structures as representative rather than assigned, and does not state which compound each code denotes.[20] Nitrogen mustard drops out of the second run for the same reason.

Across the twelve cells, three spaces by two channels by two runs, nothing survives correction (Figure 1b). Three cells reach a nominal p below 0.05, all of them in the second run: the structural space in the ultraviolet channel at 0.0167, the merged space in the same channel at 0.0194, and the electronic space in the reflectance channel at 0.0444. Benjamini–Hochberg correction over the twelve leaves a smallest adjusted value of 0.1167. The pattern of the three is itself uninformative about transfer, because the two ultraviolet cells implicate the structural and merged spaces while the reflectance cell implicates the electronic space, and the same space does not reproduce across channels of the same plate.

There is a further asymmetry in where those three cells sit, and it matters more than their size. The second run contains seven authentic agents and no designated simulant. Whatever agreement those three cells carry is agreement between one agent and another, measured on one plate, and it is not evidence that a simulant reproduces an agent. The only run that contains a designated simulant beside the agents is the first, and the first run gives nothing in any space or channel, with $|\rho| \le 0.225$ and every p at or above 0.229. The one run that could speak to simulant fidelity is silent, and the run that speaks is not about simulants.

Combining runs does not change the answer. The first and second runs share neither compounds nor plates, so their exact Mantel p values are independent and can be combined by Fisher's method. We report all six combinations of three spaces and two channels rather than selecting one, which would reintroduce the analyst freedom the test is meant to remove. The smallest adjusted value over the six is 0.2373, and none reaches 0.05. Section S13 of the Supporting Information gives the full permutation family, and Section S10 the six multiplicity

schemes under which it was read.

What this does not license is the claim that no transfer exists. Fifteen pairs is a weak test, and the pooled rank correlations are correspondingly imprecise: the structural space in the ultraviolet channel gives $\rho = +0.31$ with a 95% interval from -0.45 to $+0.81$, which excludes neither zero nor a strong association. The defensible statement is that the transfer on which simulant selection depends has not been established, in the one body of authentic-agent response data where it could be tested, and that the experiment which would settle it is a run containing a compound already measured in an existing run. Such a bridge would fix the plate correction and open the 136 cross-run pairs that are closed today.

The standard G-agent simulant against the G agents

The transfer test above treats all compounds alike. One compound deserves to be named, because the field treats it as the G-agent surrogate and because decontamination efficiency is quantified against it. The first run contains DMMP together with the four G agents and VX, so its response can be set against theirs directly.

In the ultraviolet channel DMMP correlates with tabun at $+0.764$ and with sarin at $+0.645$, and anti-correlates with soman at -0.492, cyclosarin at -0.721 and VX at -0.702. In the reflectance channel the same comparison gives -0.112, $+0.569$, -0.479, -0.495 and -0.273. The mean over all four G agents is $+0.049$ in the ultraviolet channel and -0.129 in reflectance; over the three G agents other than tabun it is -0.189 and -0.135. Neither mean carries the claim on its own, because the sign of the correlation changes with the agent and with the channel: DMMP correlates with tabun and sarin in the ultraviolet and anti-correlates with soman and cyclosarin, and tabun changes sign between the channels while sarin does not. What this measurement does not show is a compound that behaves like the G series as a class.

A weaker claim survives in one channel. If DMMP is G-like at all it should sit closer to the G series than to VX, and that contrast, taken against the mean over the three matched G

agents and excluding tabun, is +0.513 in the ultraviolet channel and +0.137 in the reflectance channel. The intervals quoted with them, +0.131 to +0.722 and −0.144 to +0.290, come from resampling the four design concentrations with replacement. Four fixed design levels are not an exchangeable sample, and the response magnitude rises monotonically with concentration, so we read these as descriptive ranges over the concentrations measured and not as tests: the ultraviolet range lies above zero and the reflectance range spans it. One optical channel is consistent with the ordering and the other does not resolve it.

Concentration changes the picture (Figure 1c). The DMMP-to-sarin cosine rises monotonically from +0.363 at 10 μM to +0.771 at 500 μM, and the DMMP-to-cyclosarin cosine falls monotonically from −0.462 to −0.700 over the same range. Monotonicity is perfect in three of the ten agent-channel series. It is also the strongest result the design can give: with four concentrations the two-sided exact permutation has a floor of $p = 0.083$, three series reach exactly that floor, and Benjamini–Hochberg correction over the ten leaves a smallest adjusted value of 0.278. We report the direction and not a significance claim.

Two limits sit on that direction. The concentrations are of analyte in a dye solution, not of vapour at an inlet, so we do not translate the range into the operating band of a point detector. And the mean response magnitude rises with concentration at a Spearman correlation of exactly +1.000 in every one of the six run-channel cells, so a similarity that grows with concentration cannot be separated here from one that grows with signal-to-noise ratio. Distinguishing the two needs a repeat-level noise estimate, which the deposited tables do not carry. What the direction supports is the narrower statement that the resemblance between this simulant and this agent is not constant over the range at which the array was characterised.

Two further readings suggested themselves to us and did not survive their own controls. We report both, because the alternative is to present artefacts as results.

The first was a ranking. DMMP is the nearest of the five candidates to tabun in the ultraviolet channel and to sarin in the reflectance channel, and it is nearest to nothing else,

and the two channels never agree on which agent it is nearest to. That reads as evidence that the choice of optical channel decides which agent the standard simulant stands in for. It is not. With five candidates the probability of rank one is 1/5 under random assignment, the expected number of first places over ten agent-channel cells is 2.0, and the observed number is two, with a one-sided binomial p of 0.62. The ten cells share DMMP's two response vectors, one per channel, so they are not ten independent trials and the binomial is an approximation; it is reported because the observed count sits on the expectation, where no correction could move the reading. The ranking carries no information.

The second was a channel decomposition. The cosine is a sum over dye channels, and the DMMP-to-tabun value of +0.764 turns out to rest on two of the 29 channels, with the largest single channel carrying 44% of the positive mass. Read alone, that makes the resemblance look fragile. Two controls show it is not a property of DMMP. Agent-to-agent pairs in the same run concentrate identically, with median two channels in the ultraviolet and one in the reflectance channel, and a null built by permuting the dye labels of the DMMP vector gives a mean of 2.00 and 2.26. Concentration of this kind is what the product of two centred unit vectors in 29 dimensions does, and it says nothing about this compound.

What remains after both controls is the set of measured values themselves, and one structural fact. The first run contains no candidate simulant other than DMMP, so the designation cannot be contrasted against an alternative in response space at all. The permutation null that the descriptor analysis of the next section relies on is, in response space, empty.

Descriptor space cannot supply the missing validation

If response space cannot establish the transfer, the question becomes whether descriptor space can at least establish the premise it is used for, that the incumbents sit closer to the agents than other compounds do. This can be asked exactly. Designating three compounds out of the non-agents in a panel is a choice among a finite set, so the incumbent designation

can be referred to every alternative designation rather than to a sampled null. The primary panel of thirty-five compounds gives 2024 designations and the reduced panel of twenty-eight gives 680.

On the primary panel the answer is flat (Figure 2). No agent of the eleven has a valid match, the smallest nominal p over the eighteen tests is 0.1917, and the Benjamini–Hochberg adjusted value is 0.6299.

What produces that zero is worth stating, because it is not chemical background. The criterion has two parts, a normalised distance of 2 or less and a distance smaller than that to every compound outside the designated set, and on this panel it is the second part that fails. Sarin, soman and cyclosarin all sit inside the threshold of DMMP, at 1.1805, 1.4583 and 1.6041. Each of them fails because diisopropyl fluorophosphate is nearer still, and tabun, which is outside the threshold at 2.5320, has diethyl cyanophosphonate nearer than DMMP. Those two compounds are not interferents. They are compounds the sensing literature itself proposes as simulants. If the comparison set is narrowed to the sixteen interferents, which is how the second part is implemented in the deposited distance table, the same three G agents pass and the count on this panel is three rather than zero. We report the wider reading as the primary one because it is the criterion stated in Section , and we report both counts because the difference between them is not a rounding matter. The designated simulants are not failing to stand out from chemical background on this panel. They are being displaced by alternatives the literature has already put forward.

Nothing in that panel supports the designation of the incumbents in particular, and the reason it does not is that better candidates exist in descriptor space.

The reduced panel, which omits seven compounds, five of them named as simulants in the literature and two of them pesticides, reads differently. Three agents acquire a valid match, and the merged space returns a nominal one-sided p of 0.0250 with the structural space at 0.0279. Whether that survives depends on which correction is applied. Benjamini–Hochberg over the eighteen tests gives 0.1632, correcting within each space gives 0.0772,

and the Westfall–Young single-step procedure, which uses the shared permutations and so handles the dependence between tests exactly, gives 0.0794. Only the scheme that fixes one statistic in advance returns a value below 0.05, at 0.0419. Significance here is a property of one correction scheme on one panel construction, and we report all four rather than the one that passes.

The surviving structural value is also an ordinary draw from the choices available to an analyst. Drawing three descriptors at random from the structural block 2000 times and repeating the whole procedure places the observed value at the 22.4th percentile of that distribution. A value at the quarter mark of what random descriptor choices produce is not evidence about chemistry. The null for that percentile, the dimensionality controls and the effect of weighting the two descriptor blocks equally are in Supporting Information Sections S15, S16 and S17.

One quantity in this family does separate the incumbents from the alternatives, and it is worth stating precisely because it is easy to over-read. Across 500 random three-descriptor subspaces, the incumbent designation reaches $p < 0.05$ in 27.4% of them. Referred to the same quantity computed for all 680 designations, where the mean is 4.5% and the median 2.4%, the incumbent designation ranks seventh of 680 ($p = 0.0103$). The incumbents are therefore unusual in how often a lucky choice of three descriptors can be found for them. That is a statement about the flexibility of the comparison, not about proximity to the agents.

One analyst freedom we did not enumerate when we set the comparison up moved the answer further than any we did. The electronic descriptors throughout are GFN2-xTB, a method parameterised for geometries and energies rather than for virtual orbital energies. Recomputing all three coordinates of the electronic space at B3LYP/def2-SVPD moves it from one valid match of eleven and a one-sided p of 0.275 to five matches and 0.0368. The five are sarin, soman and tabun matched to DMMP, and sulfur mustard and cyclosarin matched to DPGME. Only one of them, sarin to DMMP, is a pairing the designation record makes.

The NIOSH programme names DMMP for GB and for HD, the National Research Council assessment names it for VX, and the patent that names DPGME gives it no agent class at all, so the flipped result does not recover the record either. We do not promote that value to a positive finding. It is a single uncorrected statistic in one space, reached by changing a method after seeing the first answer, and the two sets of calculations sit on different geometries, so the difference carries a geometry term we cannot separate from the method term. What it establishes is narrower and bears directly on this paper's argument: the space in which the incumbents did worst is the space whose answer the choice of electronic-structure method decides. The merged space does not move at all: thirty-eight coordinates, three valid matches and a one-sided p of 0.0250 under both methods, to the precision the exact enumeration of 680 designations can report. The result we carry forward is therefore not an artefact of the semiempirical method, while the three-coordinate space is a function of it. Supporting Information Section S9 gives the full comparison and the limits on it.

Panel composition decides the answer as well. Adding back the five compounds the sensing literature also treats as simulants moves the count from three valid matches to zero, because diisopropyl fluorophosphate is nearer to three G agents than any incumbent is, and 2-chloroethyl ethyl sulfide is nearer to sulfur mustard than dichloromethane is. The incumbents do not move. What moves is whether anything closer is allowed into the comparison.

So does the way the descriptors are encoded, and that freedom is the largest of the four. Four of the thirty-eight retained coordinates are binary indicators carried by exactly one compound of the twenty-eight: an arsenic count for lewisite, a phosphorus–cyanide count for tabun, a phosphorus–sulfur count for VX and a nitro count for chloropicrin. A binary column with a single one has sample variance exactly $1/n$, so standardising it puts the carrier and every non-carrier $\sqrt{n}$ apart and adds exactly n, here 28, to the squared distance between that one compound and every other. The addition says that the compound is unique in the panel, not that it is far from DMMP, and for tabun it is 47.6% of the squared distance to DMMP. Dropping those four coordinates moves tabun from 2.0038 to 1.4509 and the

count from three valid matches to four. Dropping every binary indicator carried by three compounds or fewer gives five, and dropping every coordinate with five or fewer distinct values gives six, with the merged-space one-sided p falling from 0.0250 to 0.0074. We kept the encoding that produces the smallest count, and we report the sweep because that is the only evidence that we did not select the encoding for its answer. Supporting Information Section S6 gives the column-by-column accounting.

A fifth choice is the shape block. Nine of the thirty-eight coordinates retained on the twenty-eight-compound panel, and nine of the forty-one on the primary panel, are functions of the three principal moments of inertia. Because the relations are ratios rather than linear, the correlation filter at $|r| > 0.95$ does not remove them: the largest pairwise correlation inside the block is 0.9364 on the twenty-eight-compound panel and 0.9359 on the primary panel. They are not exact functions of one another either, since each descriptor is Boltzmann averaged over conformers independently, which leaves relative discrepancies of 1.4 to 15.3% in the three identities we checked. We therefore kept the full block and recomputed everything with the seven derived coordinates dropped, as a declared sensitivity rather than a replacement. On the twenty-eight-compound panel the shape contribution to mean squared pairwise distance falls from 23.7 to 6.5%, and for the two agents whose nearest-match rule acetonitrile blocks it falls from 69.8 and 61.9% to 0.20 and 0.12%. The count of valid matches there falls from three to zero. On the primary panel it is zero under both encodings.

Five analyst choices therefore move the count of valid matches: the panel, the order in which the correlation filter is handed its columns, the conformer embedding seed, the descriptor encoding and the shape block. None of them moves the result of Section . Recomputing the twelve Mantel tests under the encodings above leaves the smallest Benjamini–Hochberg value at 0.1167, 0.4722 and 0.3889, and dropping the derived shape coordinates raises it to 0.1708 with the largest rank correlation unchanged at 0.571. None is below 0.05, and removing the single-carrier indicators makes the agreement weaker rather than stronger. We draw the distinction explicitly because the two results are not equally load-bearing. The

agreement test between descriptor distance and measured response survives every choice we enumerated. The count of valid matches survives none of them, and a quantity that five arbitrary choices can move between zero and six is not a measurement of anything. That is this paper's argument applied to its own instrument, and it is the reason we do not report a match count as a finding.

One descriptor the opening of this paper names first is absent from the panel above, and we tested it separately rather than argue around it. Volatility enters as one axis, $\log_{10}$ of measured vapour pressure, for which all twenty-eight compounds of the reduced panel have a measurement and two compounds of the primary panel have none. Two tests were fixed before the values were seen and are held outside the family of seventy-two, because folding them in would recompute every adjusted value in the paper to answer a question one test each can settle. On this axis the incumbent trio does beat chance: the pre-specified statistic gives a one-sided p of 0.0206 over all 680 designations, ranking fourteenth. The margin is 0.0015 in standardised units, and the count of valid matches returns one match with a p of 0.4647, ranking 316th of 680. That single match pairs dipropylene glycol monomethyl ether with soman on a gap of 0.0107 decades between measurements at 26 °C and 25 °C, which is smaller than the temperature mismatch that produced it. The axis does not reproduce the pairings the usage literature implies. Dichloromethane sits 0.232 decades from hydrogen cyanide and 0.514 from phosgene; dimethyl methylphosphonate sits 0.473 from sarin while an interferent sits 0.216; nothing in the panel is within 2.659 decades of VX. We do not describe any of these as a departure from a designated pairing, because for two of the three incumbents the designating record assigns no agent class at all (Section). Agreement with measured response is positive in all four evaluable cells. Combining the two runs within each optical channel and then correcting across the two channels, the reflectance channel reaches $q = 0.0430$ while the ultraviolet channel does not ($q = 0.6389$). Volatility distance therefore does track response distance in one channel. That value is corrected within the prespecified family of two channels and not against the twelve descriptor cells, so it is not on

the same footing as the descriptor values reported above; corrected jointly with them, over sixteen cells, the smallest adjusted value is 0.1556 and nothing clears 0.05. It does not follow that the null reported here is an artefact of omitting volatility: on the pairing question the volatility axis returns one valid match at $p = 0.4647$, so it does not qualify the incumbents either. Supporting Information Section S14 sets the axis out in full.

Neither the designation record nor a toxicity model fills the gap

Two places remain where the missing evidence might be found, and neither supplies it for the incumbents. We looked because the published criterion for choosing a low-toxicity simulant is a physicochemical one,[14] and a physicochemical criterion does not by itself say what measurement the compound was validated against or how toxic it is.

The first is the standard-setting record. We traced each incumbent to the documents that designate it. DMMP is identified in a NIOSH programme for respirator materials resistance testing, where the measurement is filter breakthrough and the same document states that the simulant information is not to be used as a recommendation of any simulant.[3] It is designated a second time, for a different agent and a different measurement, in a National Research Council assessment of the Man-in-Simulant Test programme, which lists it as the fabric-testing simulant for VX and gives di-isopropyl methyl phosphonate for GB.[5] That report also states the basis on which its simulants were chosen, low toxicity and similarity of certain physical properties, and records that the chemical structure may be quite different. One compound therefore stands in for VX in one document and for GB and HD in another, against two measurements, neither of which is detection. DPGME appears in a US Army patent for a training composition, where the property required of it is detectability by ion mobility spectrometry.[4] For dichloromethane we found no designating document in the public record we searched; Supporting Information Section S7 gives the sources, the date and the three limits of that search, one of which is a Test Operations Procedure we identified but could not obtain. Each of the two traceable designations is a validation against a measurement. For

DMMP those measurements are filter breakthrough and fabric penetration, neither of which is detection. For DPGME the measurement is detection, but by ion mobility spectrometry, a different transduction from the colorimetric response reanalysed here. A simulant is validated for a purpose, and that validation does not travel with the compound to a different purpose.

The second is regulatory toxicity (Figure 3). The axis on which the agents are placed here is the EPA acute exposure guideline level, and nineteen of the twenty-eight panel compounds carry one. Two of the three incumbents do not. DMMP, the compound against which decontamination efficiency is quantified, has no AEGL, and neither has DPGME; dichloromethane has one. Those two are not absent from every register. The US Department of Energy protective action criteria give DMMP 11, 120 and 730 mg m^{-3} for PAC-1 to PAC-3, derived as temporary emergency exposure limits from a rat one-hour LC_{50}, and NIOSH gives DPGME an immediately dangerous to life or health value of 600 ppm with a recommended exposure limit of 100 ppm. What is missing is not a number of any kind but a number on the axis the agents occupy: those values rest on different averaging times, different severity definitions and different derivations, so placing them beside a ten-minute AEGL-3 would compare quantities that are not the same quantity. Supporting Information Section S23 lists what each register holds for the three incumbents with its source. A model cannot close the gap on the AEGL axis either. A random forest trained on 230 AEGL chemicals reaches $R^2 = 0.458$ in cross-validation, with an out-of-fold bias of +0.001 log units, so nothing in its internal statistics warns a user off. Applied to twenty military agents held out of training it is wrong in one direction without exception, by 0.40 to 3.42 log units. Its applicability-domain check certifies fourteen of the twenty as inside the domain, and the worst of those fourteen is wrong by 2.68 log units, a factor of 482. The mean error over the twenty is +1.68 log units with Cohen's $d_z = +1.96$. The measured axis and its gaps are tabulated in Supporting Information Section S24, and the test of whether fidelity and acute toxicity are associated at all, which finds no monotonic relation across the compounds carrying both axes, is in Section S25. Section S27 reads what that implies for the incumbent's

position on the two axes, and states why we do not turn it into a recommendation.

Conclusions

Simulant validity is not a property of a compound. It is a claim about a measurement, and the practice of designating chemical warfare agent simulants records neither the measurement a compound was validated against nor the comparison in which its validity was tested. The criterion the literature does state is physicochemical alignment with the agent,[14] which is a property of the compound and not of the measurement.

Three counts fix where the evidence stands. In the three runs reanalysed here, one designated simulant was measured beside the agents, in one run of one array, and we found no second co-measurement in the literature we searched; we state that as the limit of our search rather than as a claim about every measurement ever made. Twelve exact Mantel tests and six combinations of them leave no agreement between descriptor distance and response distance below $q = 0.05$, against a design that resolves a rank correlation of 0.85 at 80% power at that corrected threshold; Supporting Information Section S19 gives what the design could have detected at each displacement, read against the corrected multiplicity threshold. On the primary panel no agent of the eleven has a valid match to an incumbent simulant, at a Benjamini–Hochberg adjusted value of 0.6299, and the compounds that displace the incumbents there are themselves literature-proposed simulants rather than interferents.

The documentary record accounts for the first of those counts rather than contradicting it. DMMP was designated for respirator filter breakthrough testing and DPGME for an ion-mobility training composition. Neither measurement is detection, and a validation does not travel with a compound to a purpose it was never tested for. The toxicity axis stops in the same place: DMMP, the compound against which decontamination efficiency is quantified, carries no AEGL, and while other registers hold values for it they are not the quantity the agents are placed on. A random forest trained on 230 AEGL chemicals is wrong in one

direction without exception on twenty held-out military agents, so a prediction cannot stand in for the missing measurement either.

Absence of evidence is not evidence of unsuitability, and we conclude nothing about whether the incumbents are poor choices. We propose no replacement compound either, because selecting one on descriptor similarity would repeat the error this paper is about.

Two readings this paper does not support should be ruled out, because the designating documents are cited here for what they measured. DMMP is the benchmark compound against which decontamination efficiency is quantified, and nothing in this work bears on whether it is suitable for that. Filter breakthrough and surface decontamination are mass-transport and reactivity measurements, our comparison is between a colorimetric response and a descriptor vector, and a compound can be a sound surrogate for one and an unsound surrogate for the other. The same holds for DPGME and ion mobility spectrometry: the patent that designates it requires detectability by that technique, we did not measure ion mobility, and our result says nothing about whether the compound serves that purpose. The claim at issue is narrower than either. It is that a validation performed against one measurement is being carried, without a second validation, to colorimetric and spectroscopic detection of the agents themselves.

What can be specified is the measurement that would settle the question. It is a bridging run: candidate simulants and the agents they stand in for exposed on the same plate, in the same optical channels, at the same concentrations, within one session. Such a run yields the one quantity no amount of descriptor work can supply, a response-space distance between a simulant and its agent measured under identical conditions, and it needs no new chemistry and no new instrument. It needs only that the simulant be placed beside the agent rather than in place of it. Until that run exists, a statement that a compound is a validated detection simulant reports a convention, and a convention should be declared as one.

Experimental Section

Array measurements

The response data are the colorimetric tables of three published runs of one array, reanalysed here and not remeasured. The first covers tabun, sarin, soman, cyclosarin, VX and dimethyl methylphosphonate;[18] the second hydrogen cyanide, cyanogen chloride, phosgene, chloropicrin, sulfur mustard, nitrogen mustard and lewisite;[19] the third the Novichok agents A-230, A-232, A-234 and A-242.[20] We made no measurement of our own and handled no agent; what follows is what the three sources state. Each reports that the work was carried out at an Organisation for the Prohibition of Chemical Weapons designated laboratory under Chemical Weapons Convention procedures. The first names the CBRN Defense Research Institute, Seoul, as the source of the agents;[18] the third states that waste was neutralised by validated alkaline hydrolysis before disposal and does not say where its agents came from.[20] We do not describe the agents as produced at that institute, because none of the three says so. For the second run, purity was confirmed by GC–MS and exceeded 95% for every compound except lewisite, at 67%.[19] The first and third runs report no purity value for any compound, which we confirmed by reading the methods of both; the reader should treat the purity of those compounds as unstated rather than as high. Dropping lewisite changes no count reported here, on either panel and in either optical channel of the second run (Supporting Information Section S8).

The array is the same 29 commercial chromogenic and fluorogenic dyes in the same order in all three runs, commercially sourced and used as received. Stock solutions at 100 mM were prepared in ethanol, methanol, dimethyl sulfoxide, water, tetrahydrofuran or toluene and sonicated to dissolution, then diluted into acetonitrile to a working concentration of 100 μM, which holds the residual stock solvent below 0.1% (v/v) and places every dye in one solvent environment. Each analyte was delivered in acetonitrile by serial dilution to final concentrations of 0, 10, 50, 100 and 500 μM in a 2.0 mL dye solution. Samples were mixed

without incubation and imaged at once, under visible light and under ultraviolet excitation at 365 nm, with a smartphone camera fixed in an enclosed imaging box at 22 ± 3 °C and 35 ± 5% relative humidity. Regions of interest were central 20 × 10 pixel zones; RGB values were read at five positions within each region, averaged, converted to CIE $L^*a^*b^*$, and expressed as the CIEDE2000 colour difference ΔE_{00} against the 0 μM image of the same dye in the same channel. The self-referencing removes illumination and device offsets common to an image pair.

The three runs are different compounds measured on different plates and are not replicates of one another, so no ΔE_{00} value is compared across runs; every analysis below is run within a run and channel, and results are combined only at the level of correlation coefficients. The two channels are physically different measurements rather than repeats and are treated as independent modalities. The tables hold one value per compound, concentration, dye and channel.

Response patterns

Each exposure gives a 29-dye ΔE_{00} vector. Vectors were centred dye-wise within each concentration and normalised to unit length, so that agreement is measured between patterns and not between magnitudes, and agreement was taken as the cosine. Dye columns that are bit-identical across every row carry no information and were counted once, which leaves 23 effective channels in the ultraviolet pivot of the second run and 25 in that of the third.

Descriptors

Structural and physicochemical descriptors, 51 after excluding run metadata, were computed with RDKit from canonical SMILES. Ten are read from a three-dimensional structure and are taken as Boltzmann-weighted means over an ETKDGv3 conformer ensemble at 298 K with MMFF94 energies, or UFF for lewisite 1, the one compound without MMFF parameters, so that no result depends on an undeclared choice of conformer; Supporting Information

Section S4 reports what the two treatments do to the results. Electronic descriptors were computed with GFN2-xTB in the gas phase, Boltzmann averaged over up to five conformers at 298.15 K. The electronic space carries E_{HOMO}, E_{LUMO} and the dipole moment. The Parr electrophilicity index $\omega = \mu^2/2\eta$ is reported but admitted to no distance, because it is an exact function of the two orbital energies that accompany it and a filter built on linear correlation does not detect that. It and the condensed Fukui index at phosphorus, which exists only for the six phosphorus-bearing compounds, are reported in Supporting Information Section S20 and admitted to no distance. Dipole moments were recomputed at B3LYP/def2-SVPD, together with the two frontier orbital energies. The merged space, which carries the result reported above, is unchanged by that substitution. The three-coordinate electronic space is not, and we report what it does in Section rather than leaving the check described as a confirmation (Supporting Information Section S9). Identifiers and formulae were resolved against PubChem, and no property value was taken from a prediction service.

Distances and tests

Within each space, descriptors of zero variance were removed, then one of each pair correlated at $|r| > 0.95$, retaining a protected set. Remaining descriptors were standardised and distances taken as Euclidean, normalised by the median nearest-neighbour separation of the same space, which makes a normalised distance comparable within a panel and not between panels.

A paired comparison of the distance to the nearest incumbent simulant against the distance to the nearest interferent is confounded by group size, since the minimum of many draws is smaller than the minimum of few by construction. We therefore enumerated every way of designating three of the non-agents as simulants and report the rank of the actual designation as an exact p-value: $\binom{24}{3} = 2024$ designations on the primary panel and $\binom{17}{3} = 680$ on the twenty-eight-compound panel. The hypothesis is directional, so the reported p is one-sided. We give the two-sided values here rather than only the convention, because on

the twenty-eight-compound panel they decide the answer: the structural space moves from 0.0279 one-sided to 0.0559 at twice the smaller tail and 0.0441 scored as distance from the null mean, and the merged space from 0.0250 to 0.0500 and 0.0279. Under the first of the two-sided conventions neither space reaches 0.05, and the one correction scheme that returns a value below 0.05 does so only on the one-sided reading. The Mantel tests of Section are two-sided throughout, since the direction of a matrix correlation is not fixed in advance by the hypothesis. A valid match requires a normalised distance of 2 or less together with a distance smaller than that to every compound outside the designated set. Sensitivity to descriptor choice was assessed by resampling the merged-space descriptors with replacement 1000 times under a fixed seed, restandardising and recomputing within each replicate; the intervals so obtained are percentile ranges over descriptor choice and not confidence intervals in the sampling sense.

Agreement between descriptor distance and response distance was tested by Mantel correlation with exact enumeration of relabellings in each run-channel cell, with the independent runs combined by Fisher's method and the family of tests corrected by the Benjamini–Hochberg procedure. Six distance statistics and six multiplicity schemes were carried through in parallel so that the conclusion can be read against the analyst choices rather than from one of them.

Software

NumPy 2.5.3,[21] SciPy 1.18.1,[22] pandas 3.0.5,[23] scikit-learn 1.9.1,[24] Matplotlib 3.11.2,[25] RDKit 2026.03.6[26] and PySCF 2.14.0.[27] Conformer embedding used ETKDGv3,[28,29] the semiempirical electronic descriptors GFN2-xTB,[30,31] and the physicochemical descriptors the atomic-contribution and fragment methods implemented in RDKit.[32,33] All analyses are deterministic given the seed. The designation tests on both panels and the Mantel tests are exhaustive; the eight-simulant designation null and the random-subspace controls are sampled, at 20000 and 2000 draws from the recorded seed. Analysis code, descriptor tables, the

parsed regulatory values with their sources and every figure script are deposited in a public repository,[34] from which every number we computed can be regenerated. Three classes of number are quoted rather than recomputed and are marked as such where they appear: values taken from the three source studies, values taken from regulatory registers, and the two cross-validated classification accuracies that the deposited tables do not reproduce. The colorimetric input is the deposited ΔE_{00} table; the conversion from the original plate images to those colour differences was performed in the source studies and is outside this deposit.

Acknowledgement

The colorimetric response data reanalysed here were measured in three earlier studies from this laboratory,[18–20] and we thank the staff of the CBRN Defense Research Institute who performed those measurements under Chemical Weapons Convention procedures. K.K. carried out the conceptualisation, software, formal analysis and visualisation and wrote the original draft. J.K. and J.Y. contributed to the analysis and validation, J.Y. to the original draft, M.K. to the validation and to reviewing and editing, D.-H.L. to validation, resources and supervision, and Y.J.J. to validation, supervision and project administration. This research received no specific grant from any funding agency in the public, commercial or not-for-profit sectors. The authors declare no competing financial interest.

Supporting Information Available

The Supporting Information is available free of charge at `https://pubs.acs.org/doi/XX.XXXX/acssensors.XXXXXXX`.

- Acute inhalation toxicity of the panel compounds and the gaps in it; predictive interpolation of that axis and why a model cannot fill the gaps; the decoupling test between fidelity and acute toxicity (Sections S3 and S22); conformer dependence of

the shape descriptors and the reproduction gate; the density-functional check on the electrophilicity axis and on the dipole; the full permutation family under all six multiplicity schemes; dimensionality, subspace and descriptor-block weighting controls, each with its null; the measured colorimetric array response and what it does and does not establish; phosphorus-local reactivity for the six phosphorus-bearing compounds; and the structural filter (Section S26) (PDF)

References

(1) Askim, J. R.; Mahmoudi, M.; Suslick, K. S. Optical sensor arrays for chemical sensing: the optoelectronic nose. *Chemical Society Reviews* **2013**, *42*, 8649.

(2) Rakow, N. A.; Suslick, K. S. A colorimetric sensor array for odour visualization. *Nature* **2000**, *406*, 710–713.

(3) National Institute for Occupational Safety and Health *Chemical Warfare Agent (CWA) Simulant Program for Respirator Materials Resistance Testing*; 2002; Dated 31 July 2002.

(4) Seitzinger, A. T.; Genovese, J. A. Chemical Agent Simulant Training Composition. 2006; `https://patents.google.com/patent/US7129094B1/en`, U.S. Patent US 7,129,094 B1, assigned to the United States Department of the Army. Filed 16 May 2003; granted 31 October 2006.

(5) National Research Council *Technical Assessment of the Man-in-Simulant Test Program*; National Academies Press: Washington, DC, 1997; Chapter 3, Simulant Selection.

(6) ASTM International Standard Test Method for Man-In-Simulant Test (MIST) for Protective Ensembles. ASTM F2588-12(2020), 2020; Designation F2588-12, reapproved 2020.

(7) Lavoie, J.; Srinivasan, S.; Nagarajan, R. Using cheminformatics to find simulants for chemical warfare agents. *Journal of Hazardous Materials* **2011**, *194*, 85–91.

(8) Bartelt-Hunt, S. L.; Knappe, D. R. U.; Barlaz, M. A. A Review of Chemical Warfare Agent Simulants for the Study of Environmental Behavior. *Critical Reviews in Environmental Science and Technology* **2008**, *38*, 112–136.

(9) James, T.; Wyke, S.; Marczylo, T.; Collins, S.; Gaulton, T.; Foxall, K.; Amlôt, R.; Duarte-Davidson, R. Chemical warfare agent simulants for human volunteer trials of emergency decontamination: A systematic review. *Journal of Applied Toxicology* **2018**, *38*, 113–121.

(10) James, T.; Collins, S.; Marczylo, T. Identification of novel simulants for toxic industrial chemicals and chemical warfare agents for human decontamination studies: A systematic review and categorisation of physicochemical characteristics. *International Journal of Environmental Research and Public Health* **2021**, *18*, 8681.

(11) US Army Test and Evaluation Command *Test Operations Procedure (TOP) 08-2-196: Simulant Selection for Laboratory, Chamber, and Field Testing*; 2011; Cited in TOP 08-2-188 CN1 as the governing simulant-selection procedure; the document itself was not obtainable for this work.

(12) Agrawal, M.; Sava Gallis, D. F.; Greathouse, J. A.; Sholl, D. S. How Useful Are Common Simulants of Chemical Warfare Agents at Predicting Adsorption Behavior? *The Journal of Physical Chemistry C* **2018**, *122*, 26061–26069.

(13) Lawrence, D. S.; Van Houten, K. A. Simulants for Chemical and Explosive Threats. *Johns Hopkins APL Technical Digest* **2025**, *37*.

(14) Webb, S.; Coulon, F.; Temple, T. A critical review of liquid, low toxicity chemical warfare agent simulants: Enhancing accuracy, safety, and methodological approaches for sampling. *Journal of Hazardous Materials* **2025**, *492*, 138021.

(15) Chen, Y.; Liu, J. A Practical Guide of Choosing and Validating Aptamers for Biosensors. *ACS Sensors* **2026**, Advance online publication.

(16) Yoo, J.; Lee, D.; Lee, S.; Kang, S.; Kim, H. I.; Jang, Y. J.; Kim, J.; Park, T. H. Combinatorial Pattern Response of Bioelectronic Nose for the Detection of Real Nerve Agents. *ACS Sensors* **2025**, *10*, 185–195.

(17) Bae, S.; Kang, K.; Kim, Y. K.; Jang, Y. J.; Lee, D.-H. Field-deployable real-time AI System for chemical warfare agent detection using YOLOv8 and colorimetric sensors. *Chemometrics and Intelligent Laboratory Systems* **2025**, *261*, 105365.

(18) Kim, J.; Kang, K.; Shin, M.; Kim, S.; Yoo, J.; Kim, M.-K.; Lee, W. B.; Lee, D.-H. Machine learning-enhanced colorimetric sensor array for rapid detection of nerve agents. *Journal of Hazardous Materials* **2026**, *501*, 140703.

(19) Kim, S.; Yoo, J.; Kang, K.; Kim, J.; Shin, M.; Cha, Y. K.; Churchill, D. G.; Kim, M.-K.; Lee, D.-H. A multiplex chromogenic–fluorogenic sensor array for comprehensive optical discrimination of chemical warfare agents. *Discover Sensors* **2025**, *1*, 29.

(20) Shin, M.; Kang, K.; Kim, J.; Kim, S.; Yoo, J.; Jang, Y. J.; Kim, M.-K.; Anslyn, E. V.; Lee, D.-H. Portable machine learning–integrated optical sensor array for rapid and selective detection of authentic Novichok nerve agents. *Sensors and Actuators B: Chemical* **2026**, *457*, 139706.

(21) Harris, C. R. et al. Array programming with NumPy. *Nature* **2020**, *585*, 357–362.

(22) Virtanen, P. et al. SciPy 1.0: fundamental algorithms for scientific computing in Python. *Nature Methods* **2020**, *17*, 261–272.

(23) McKinney, W. *Data Structures for Statistical Computing in Python*; 2010; pp 56–61.

(24) Pedregosa, F. et al. Scikit-learn: Machine Learning in Python. *Journal of Machine Learning Research* **2011**, *12*, 2825–2830.

(25) Hunter, J. D. Matplotlib: A 2D Graphics Environment. *Computing in Science & Engineering* **2007**, *9*, 90–95.

(26) Landrum, G.; others RDKit: Open-source cheminformatics. `https://www.rdkit.org`, 2026; Version 2026.03.6. Cited by name, version and project URL: the Zenodo DOIs registered for RDKit identify individual releases (10.5281/zenodo.10398 is release 2014.03.1) and no version-independent concept DOI could be verified, so none is given here.

(27) Sun, Q.; Berkelbach, T. C.; Blunt, N. S.; Booth, G. H.; Guo, S.; Li, Z.; Liu, J.; McClain, J. D.; Sayfutyarova, E. R.; Sharma, S.; Wouters, S.; Chan, G. K.-L. PySCF: the Python-based simulations of chemistry framework. *WIREs Computational Molecular Science* **2018**, *8*, e1340.

(28) Riniker, S.; Landrum, G. A. Better Informed Distance Geometry: Using What We Know To Improve Conformation Generation. *Journal of Chemical Information and Modeling* **2015**, *55*, 2562–2574.

(29) Wang, S.; Witek, J.; Landrum, G. A.; Riniker, S. Improving Conformer Generation for Small Rings and Macrocycles Based on Distance Geometry and Experimental Torsional-Angle Preferences. *Journal of Chemical Information and Modeling* **2020**, *60*, 2044–2058.

(30) Grimme, S.; Bannwarth, C.; Shushkov, P. A Robust and Accurate Tight-Binding Quantum Chemical Method for Structures, Vibrational Frequencies, and Noncovalent Interactions of Large Molecular Systems Parametrized for All spd-Block Elements (Z = 1–86). *Journal of Chemical Theory and Computation* **2017**, *13*, 1989–2009.

(31) Bannwarth, C.; Ehlert, S.; Grimme, S. GFN2-xTB—An Accurate and Broadly Parametrized Self-Consistent Tight-Binding Quantum Chemical Method with Multi-

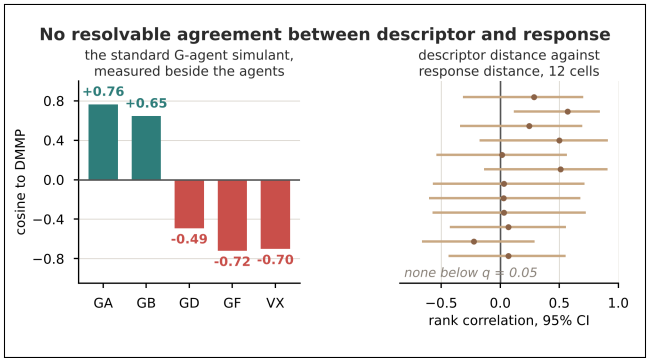
pole Electrostatics and Density-Dependent Dispersion Contributions. *Journal of Chemical Theory and Computation* **2019**, *15*, 1652–1671.

(32) Wildman, S. A.; Crippen, G. M. Prediction of Physicochemical Parameters by Atomic Contributions. *Journal of Chemical Information and Computer Sciences* **1999**, *39*, 868–873.

(33) Ertl, P.; Rohde, B.; Selzer, P. Fast Calculation of Molecular Polar Surface Area as a Sum of Fragment-Based Contributions and Its Application to the Prediction of Drug Transport Properties. *Journal of Medicinal Chemistry* **2000**, *43*, 3714–3717.

(34) Kang, K. simulant-designation-evidence: analysis code, descriptor tables and figure scripts for this work. `https://github.com/bisu9082/simulant-designation-evidence`, 2026; Public Git repository.

TOC Graphic


No resolvable agreement between descriptor and response
the standard G-agent simulant, measured beside the agents
cosine to DMMP
0.8
0.4
0.0
−0.4
−0.8
+0.76
+0.65
-0.49
-0.72
-0.70
GA
GB
GD
GF
VX
descriptor distance against response distance, 12 cells
none below q = 0.05
−0.5
0.0
0.5
1.0
rank correlation, 95% CI

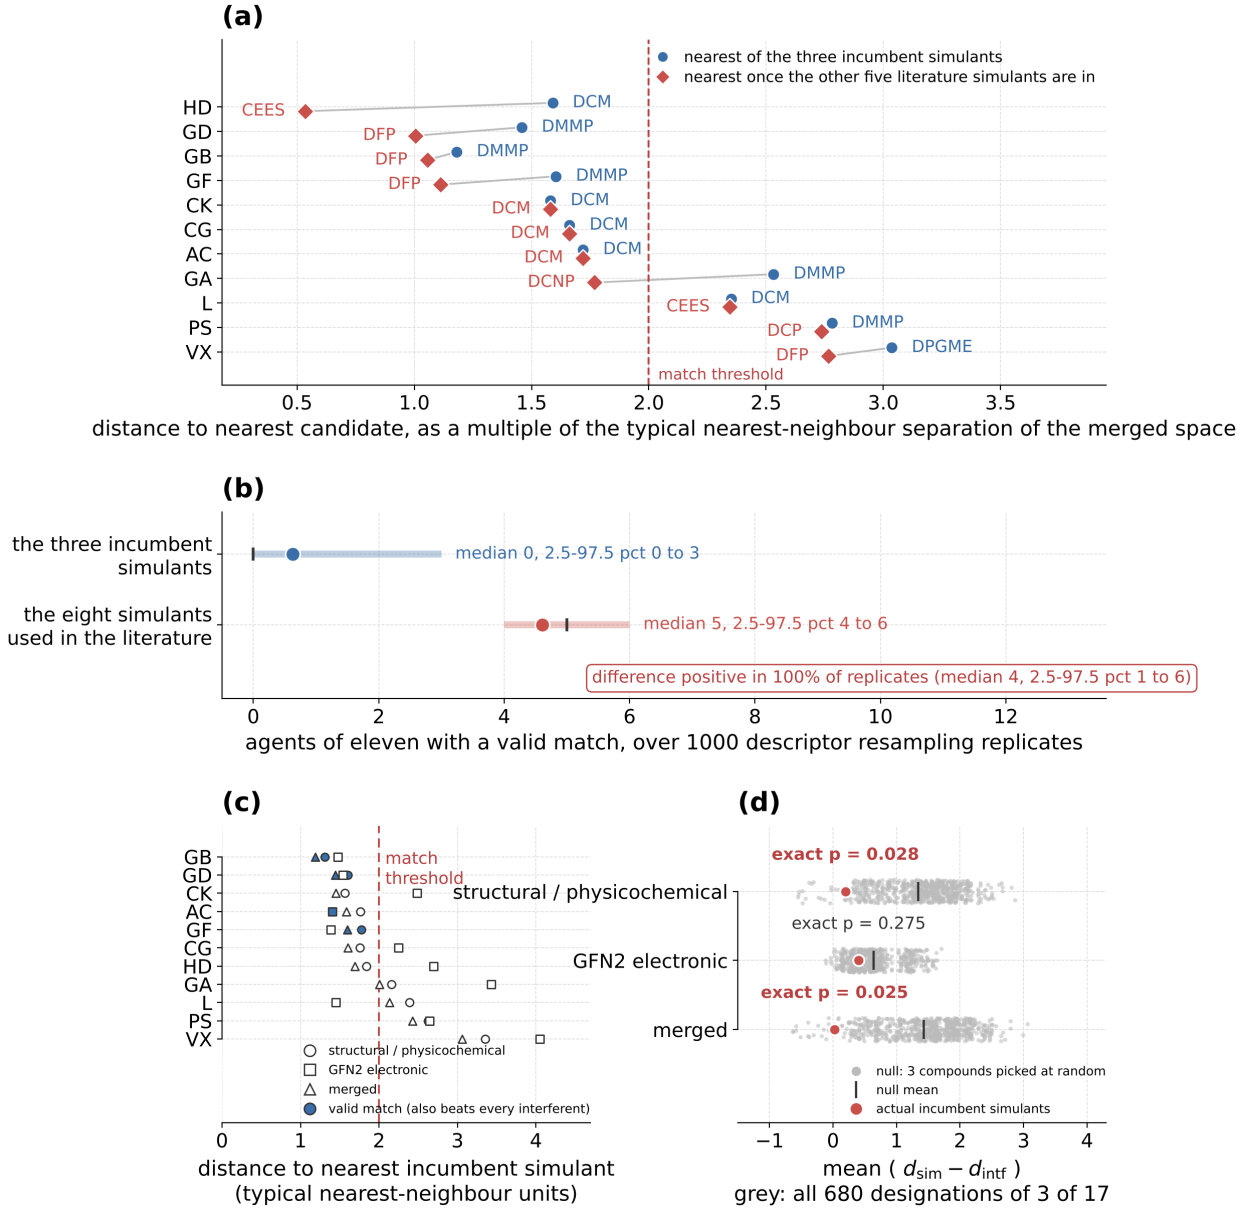


Figure 2: Descriptor-space tests. (a) Distance from each agent to its nearest candidate, for the three incumbents and for the eight compounds the sensing literature also calls simulants. (b) Valid matches over 1000 descriptor resampling replicates. (c) Distance to the nearest incumbent by space. (d) Observed mean difference against the exact null over all 680 designations of three compounds from seventeen.

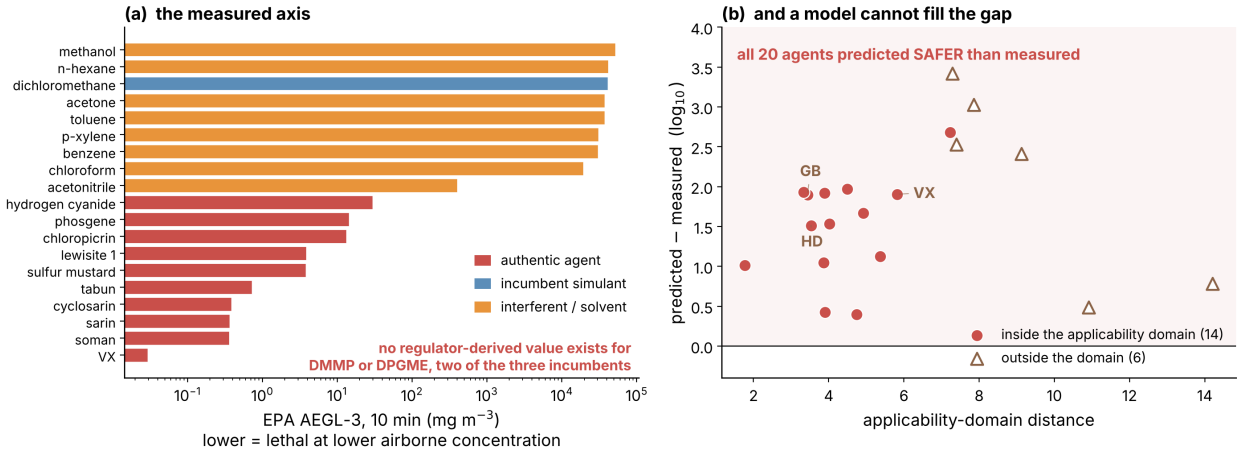


Figure 3: The toxicity axis. (a) Measured EPA AEGL-3 10-minute values for the panel compounds that carry one. (b) Prediction error of a random forest trained on 230 AEGL chemicals, for twenty military agents held out of training, against distance to the training set.